



# Machine Learning

## Session 7 - Reinforcement Learning



hadrien.salem@centralelille.fr



[introduction-to-data-science](#)

# Introduction

What did we do last time?

# Course outline

## Machine learning course

Session 1: Regression

Session 2: Supervised classification

Session 3: Clustering

Session 4: Decision trees and ensemble methods

Session 5: Introduction to neural networks

Session 6: Advanced neural networks


**Session 7:** Introduction to reinforcement learning

**Session 8:** Reading science papers




**Project**

# What is reinforcement learning?



Through trial and error, an agent keeps learning in an interactive environment ([source](#))



**Agents** take actions  
in an **environment**  
aimed at maximizing their  
**cumulative reward**

Desired behaviors are rewards, undesired behaviors are punished.  
Through trial and error, the agent learns the best behavior.



# Vocabulary

**Agent** - The decision maker and learner

**Environment** - A physical world where an agent learns and decides the actions to be performed

**Action** – Interaction of the agent with its environment


**State** – The current situation of the agent in the environment

**Reward** – Feedback (scalar value) given to the agent by the environment as a result of taking an action

**Policy** – A map of situations to actions, the strategy of the agent

**Value Function** – Function evaluating how good it is to be in a particular state or state-action pair

# Applications of RL



In 2016, AlphaGo beats professional Go player Lee Sedol



# Applications of reinforcement learning

## **Game playing**

**Robotics** - Robots learn independently how to perform certain tasks (walking, picking, placing)

**Recommendation systems** – Youtube videos, Spotify playlists, etc.

**NLP** – Chatbots learn from their past experiences

**Education** – Personalizing the students' learning experiences

**Dynamic parameter optimization** – Improve the performance of an algorithm as it is being used

# How reinforcement learning works

## The example of Q learning

Take action



Compute reward



Update Q table

Game Board:



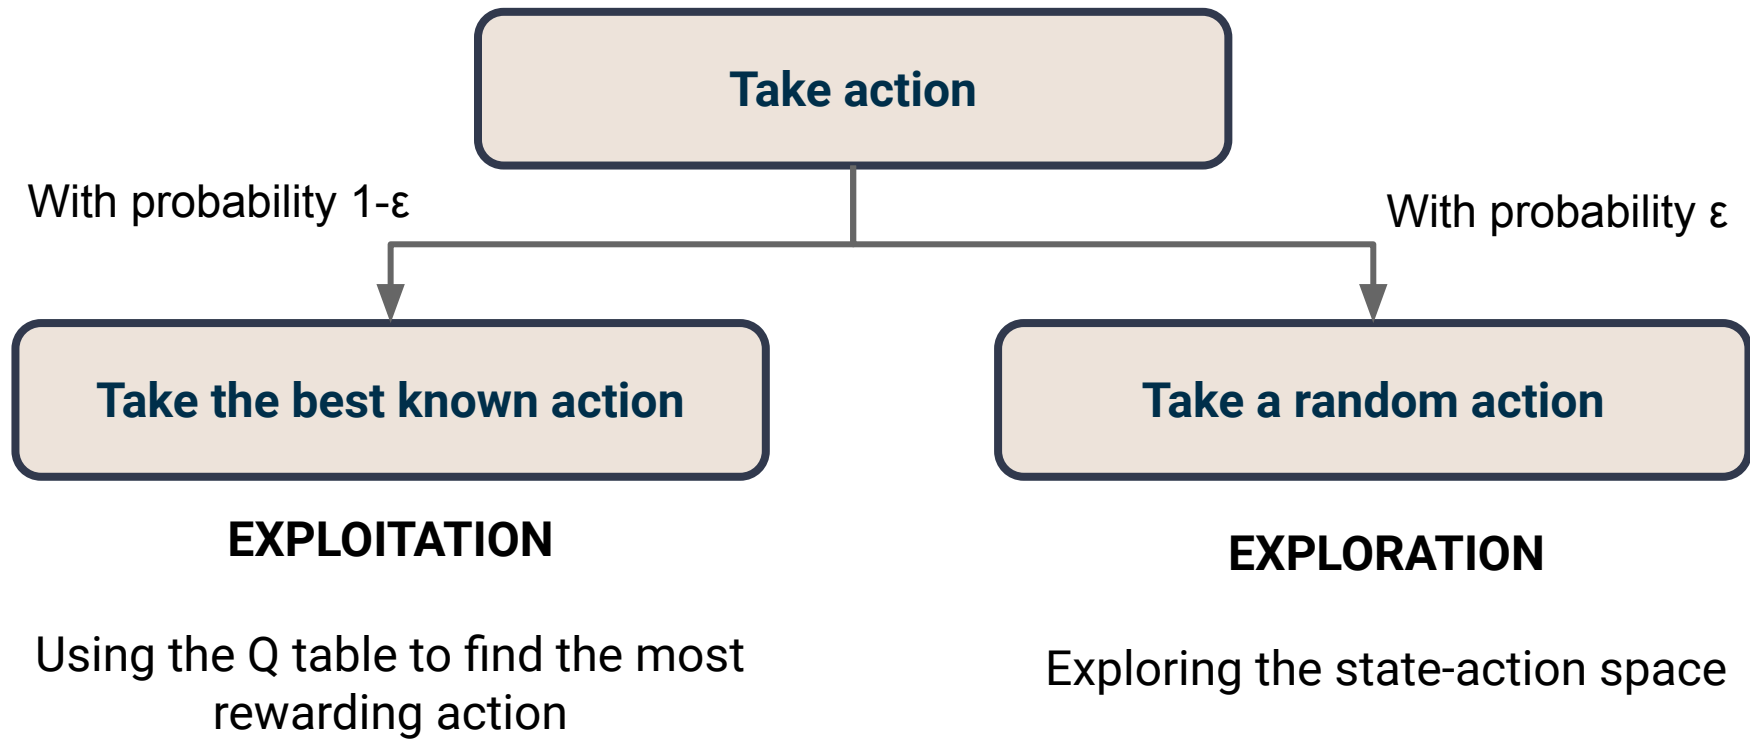
Current state (s):  
0 0 0  
0 1 0

Q Table:

$\gamma = 0.95$

	0 0 0 1 0 0	0 0 0 0 1 0	0 0 0 0 0 1	1 0 0 0 0 0	0 1 0 0 0 0	0 0 1 0 0 0
↑	0.2	0.3	1.0	-0.22	-0.3	0.0
↓	-0.5	-0.4	-0.2	-0.04	-0.02	0.0
→	0.21	0.4	-0.3	0.5	1.0	0.0
←	-0.6	-0.1	-0.1	-0.31	-0.01	0.0

With Q learning, agents base their decisions on a Q table





$$\underbrace{Q(S_t, A_t)}_{\text{New Q-value estimation}} \leftarrow \underbrace{Q(S_t, A_t)}_{\text{Former Q-value estimation}} + \underbrace{\alpha}_{\text{Learning Rate}} [\underbrace{R_{t+1}}_{\text{Immediate Reward}} + \underbrace{\gamma \max_a Q(S_{t+1}, a)}_{\text{Discounted Estimate optimal Q-value of next state}} - \underbrace{Q(S_t, A_t)}_{\text{Former Q-value estimation}}]$$

New  
Q-value  
estimation

Former  
Q-value  
estimation

Learning  
Rate

Immediate  
Reward

Discounted Estimate  
optimal Q-value  
of next state

Former  
Q-value  
estimation

TD Target

TD Error

- **TD Target (Temporal-Difference Target):** estimate of the expected cumulative reward
- **TD Error (Temporal-Difference Error):** quantifies how far off the current value estimate is from what is expected

# Variants

# Practical work

The notebook contains all the necessary instructions

# Debrief

# Debrief – G3



<https://forms.gle/TiR9uaYTFF3iTX9e6>

# Debrief

**What did we learn today?**

**What could we have done better?**

**What are we doing next time?**